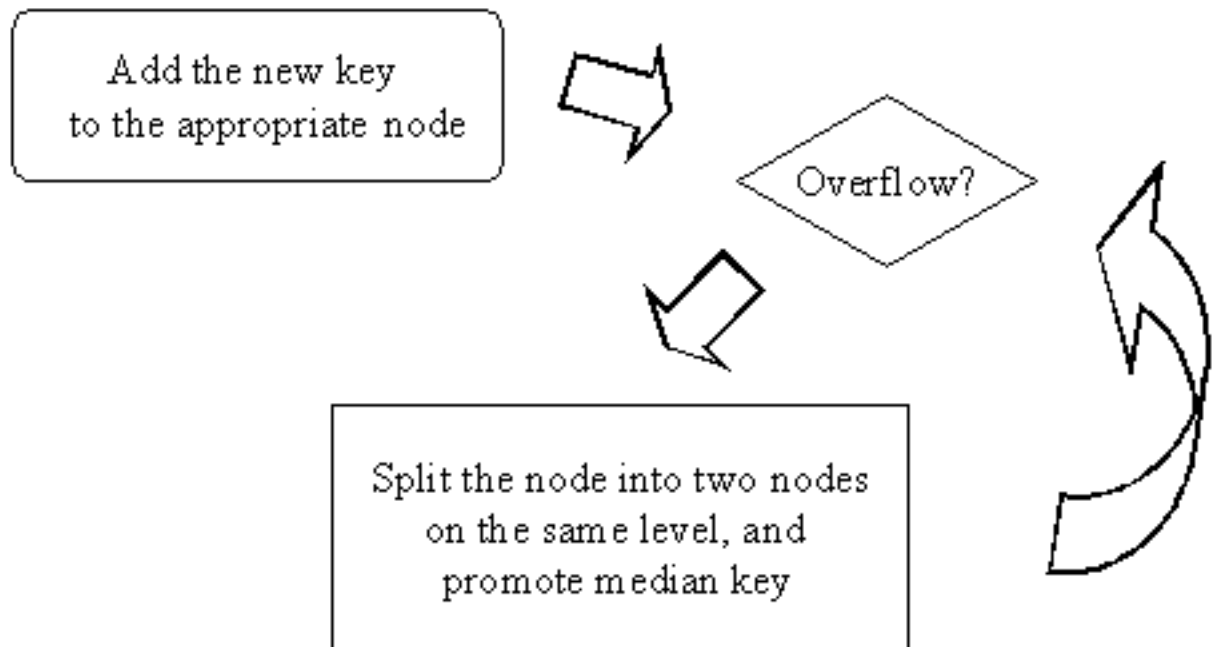


Definition of a B-Tree

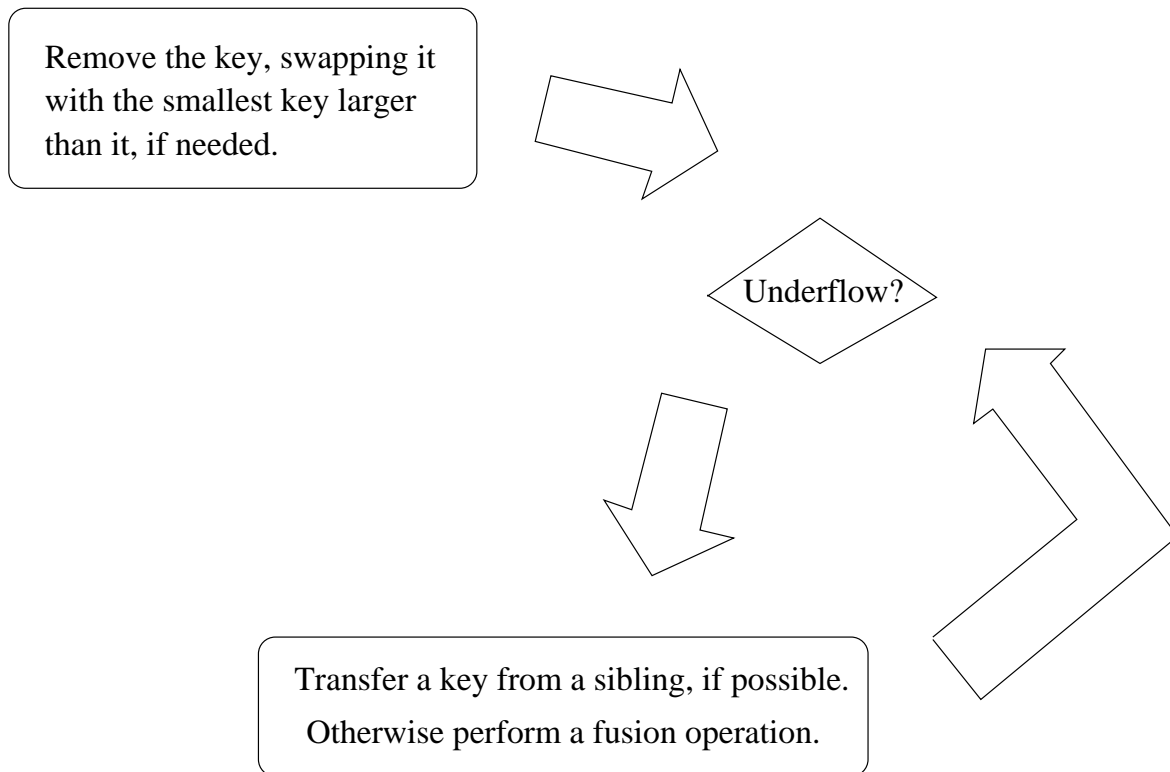
A B-tree of order d is a multi-way search tree T with the following properties:

1. The root of T has at least 2 children and at most d children.
2. All internal nodes of T (other than the root) have between $\lceil d/2 \rceil$ and d children.
3. All external nodes of T are at the same level.

Insertion in a B-Tree



Deletion from a B-Tree



External Search

- Main memory is several orders of magnitude faster than disk. Thus the main goal of maintaining a dictionary in disk is to minimize the number of disk accesses.
- To try to minimize the time needed to transfer information from disk to main memory, data items on a disk are grouped into contiguous sections called *blocks*.
- B-trees are used to implement dictionaries in external memory (disk).
- The order of a B-tree is chosen so that the size of a node is smaller than the size of a disk block, but as close as possible.

Height of a B-tree

- The height of a B-tree gives an upper bound on the maximum number of disk accesses required to access information stored in it.
- The height of a B-tree of order d is

$$O(\log_{\lceil d/2 \rceil} n),$$

where n is the number of keys stored in the tree.

- For example, the height of a B-tree of order 200 storing one million keys, is only 3.