

Chapter 14

The Preprocessor

Introduction

- Directives such as `#define` and `#include` are handled by the *preprocessor*, a piece of software that edits **C** programs just prior to compilation
- The preprocessor is a powerful tool, but it also can be a source of hard-to-find bugs

How the Preprocessor Works

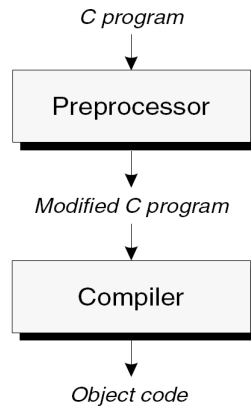
- The preprocessor looks for *preprocessing directives*, which begin with a # character
- So far, we have encountered the #define and #include directives
- #define defines a *macro*—a name that represents something else, such as a constant
- The preprocessor responds to a #define directive by “*storing*” the name of the macro along with its definition
- When the macro is used *later*, the preprocessor “*expands*” the macro, replacing it by its defined value

How the Preprocessor Works

- #include tells the preprocessor to open a particular file and “*include*” its contents as part of the file being compiled
- For example, the line
#include <stdio.h>
instructs the preprocessor to open the file named `stdio.h` and bring its contents into the program

How the Preprocessor Works

- The preprocessor's role in the compilation process:



How the Preprocessor Works

- The input to the preprocessor is a **C** program, possibly containing directives
- The preprocessor
 - **executes** these directives
 - **removing** them in the process
- The preprocessor's output goes directly into the compiler

How the Preprocessor Works

- The **celsius.c** program

```
/* Converts a Fahrenheit temperature to Celsius */
#include <stdio.h>
#define FREEZING_PT 32.0f
#define SCALE_FACTOR (5.0f / 9.0f)

int main(void)
{
    float fahrenheit, celsius;

    printf("Enter Fahrenheit temperature: ");
    scanf("%f", &fahrenheit);

    celsius = (fahrenheit - FREEZING_PT) * SCALE_FACTOR;
    printf("Celsius equivalent is: %.1f\n", celsius);

    return 0;
}
```

How the Preprocessor Works

- The program after preprocessing:

```
Blank line
Blank line
Lines brought in from stdio.h
Blank line
Blank line
Blank line
Blank line
int main(void)
{
    float fahrenheit, celsius;

    printf("Enter Fahrenheit temperature: ");
    scanf("%f", &fahrenheit);

    celsius = (fahrenheit - 32.0f) * (5.0f / 9.0f);
    printf("Celsius equivalent is: %.1f\n", celsius);

    return 0;
}
```

How the Preprocessor Works

- The preprocessor does a bit more than just execute directives
- In particular,
 - it *replaces* each comment with a *single* space character
- Some preprocessors go further and:
 - remove unnecessary white-space characters, including spaces and tabs at the beginning of indented lines
- In the early days of **C**, the preprocessor was a separate program

How the Preprocessor Works

- Most **C** compilers provide a way to view the output of the preprocessor
- Some compilers generate preprocessor output when a certain option is specified (**gcc** and **cc** will do so when the **-E** option is used)

Preprocessing Directives

- Most preprocessing directives fall into one of three categories:
 - **Macro definition**
 - The `#define` directive defines a macro
 - The `#undef` directive removes a macro definition
 - **File inclusion**
 - The `#include` directive causes the contents of a specified file to be included in a program
 - **Conditional compilation**
 - The `#if`, `#ifdef`, `#ifndef`, `#elif`, `#else`, and `#endif` directives allow blocks of text to be either *included in* or *excluded from* a program

Preprocessing Directives

Several rules apply to all directives

- **Directives always begin with the # symbol**
The `#` symbol *need not* be at the beginning of a line, as long as *only white space precedes it*
- **Any number of spaces and horizontal tab characters may separate the tokens in a directive**

Example:

```
#      define      N      100
```

Preprocessing Directives

- ***Directives always end at the first new-line character, unless explicitly continued***

To continue a directive to the next line, end the current line with a \ character (i.e., \ character immediately followed by a ***new-line***):

```
#define DISK_CAPACITY (SIDES *  
                        TRACKS_PER_SIDE *  
                        SECTORS_PER_TRACK *  
                        BYTES_PER_SECTOR)
```

Preprocessing Directives

- ***Directives can appear any where in a program***
Although #define and #include directives *usually* appear at the beginning of a file, other directives are more likely to show up later
- ***Comments may appear on the same line as a directive***

It is *good practice* to put a comment at the end of a macro definition:

```
#define FREEZING_PT 32.0f /* freezing point of water */
```

Macro Definitions

- The macros that we have been using since Chapter 2 are known as *simple* macros, because they have no parameters
- The preprocessor also supports *parameterized* macros

Simple Macros

- Definition of a *simple macro* (or *object-like macro*):
`#define identifier replacement-list`
- *replacement-list* is any sequence of *preprocessing tokens*
- It is *legal* to have *empty replacement-list*
- The replacement list may include:
identifiers, keywords, numeric constants,
punctuation, operators, character constants,
and string literals
- Wherever the *identifier* appears *later* in the file, the preprocessor substitutes it with the *replacement-list*

Simple Macros

- When macros are used as constants, **C** programmers usually capitalize all letters in their names
- However, there is no consensus as to how to capitalize macros used for other purposes
 - Some programmers like to draw attention to macros by using all upper-case letters in their names
 - Others prefer lower-case names

Simple Macros

- Any extra symbols in a macro definition will become part of the replacement list
- Putting the = symbol in a macro definition is a common error:

```
#define N = 100  /** WRONG **/  
...  
int a[N];        /* becomes int a[= 100]; */
```
- Ending a macro definition with a semicolon is another popular mistake:

```
#define N 100;  /** WRONG **/  
...  
int a[N];        /* becomes int a[100;]; */
```
- The compiler will detect most errors caused by extra symbols in a macro definition

Simple Macros

- Simple macros are primarily used for defining “*manifest constants*”, i.e., names that represent numeric, character, and string values:

```
#define STR_LEN 80
#define TRUE 1
#define FALSE 0
#define PI 3.14159
#define CR '\r'
#define EOS '\0'
#define MEM_ERR "Error: not enough memory"
```

Simple Macros

- Advantages of using `#define` to create names for constants:
 - *It makes programs easier to read*
The name of the macro can help the reader understand the meaning of the constant
 - *It makes programs easier to modify*
We can change the value of a constant throughout a program by modifying a single macro definition
 - *It helps avoid inconsistencies and typographical errors*
If a numerical constant like 3.14159 appears many times in a program, chances are it will occasionally be written 3.1416 or 3.14195 by accident

Simple Macros

- Simple macros have additional uses
- ***Making minor changes to the syntax of C***

Macros can serve as alternate names for **C** symbols:

```
#define BEGIN {  
#define END   }  
#define LOOP for(;;)
```

Changing the syntax of **C** usually *is not a good idea*, since it can make programs harder for others to understand

Simple Macros

- ***Renaming types***

An example from Chapter 5:

```
#define BOOL int
```

Type definitions are a better alternative

- ***Controlling conditional compilation***

Macros play an important role in controlling conditional compilation

A macro that might indicate “*debugging mode*”:

```
#define DEBUG
```

Parameterized Macros

- Definition of a **parameterized macro** (also known as a **function-like macro**):
`#define identifier(x1 , x2 , ... , xn) replacement-list`
 x_1, x_2, \dots, x_n are identifiers (the macro's **parameters**)
- The parameters may appear as many times as desired in the replacement list
- There must be **no space** between the **macro name** and the **left parenthesis**
- **If a space is left between the macro name and the left parenthesis, the preprocessor will treat (x₁, x₂, ..., x_n) as part of the replacement list**

Parameterized Macros

- Wherever a macro **invocation** of the form
`identifier(y1, y2, ..., yn)` appears later in the program,
the preprocessor replaces it with **replacement-list**,
substituting y_1 for x_1 , y_2 for x_2 , and so forth
- Parameterized macros often serve as simple functions

Parameterized Macros

- Examples of parameterized macros:

```
#define MAX(x,y)    ((x)>(y)?(x):(y))
#define IS_EVEN(n) ((n)%2==0)
```

- Invocations of these macros:

```
i = MAX(j + k, m - n);
if (IS_EVEN(i)) i++;
```

- The same lines after macro replacement:

```
i = ((j + k)>(m - n)?(j + k):(m - n));
if (((i)%2==0)) i++;
```

Parameterized Macros

- A more complicated function-like macro:

```
#define TOUPPER(c) \
    ((c)>='a'&&(c)<='z'? (c) - 'a' + 'A' : (c))
```

- The `<ctype.h>` header provides a similar function named `toupper`
- A parameterized macro may have an empty parameter list:

```
#define getchar() getc(stdin)
```

- The empty parameter list is not really needed, but it makes `getchar` resemble a function

Parameterized Macros

- Using a parameterized macro instead of a true function has a couple of advantages:
 - ***The program may be slightly faster***
A function call usually requires some overhead during program execution, but a macro invocation does not
 - ***Macros are “generic”***
A macro can accept arguments of any type, provided that the resulting program is valid

Parameterized Macros

- Parameterized macros also have disadvantages
 - ***The compiled code will often be larger***
Each macro invocation increases the size of the source program (and hence the compiled code)
The problem is compounded when macro invocations are nested:

```
n = MAX(i, MAX(j, k));
```


The statement after preprocessing:

```
n = ((i) > ((j) > (k) ? (j) : (k))) ? (i) : (((j) > (k) ? (j) : (k))) ;
```

Parameterized Macros

– *Arguments are not type-checked*

When a function is called, the compiler checks each argument to see if it has the appropriate type

Macro arguments are not checked by the preprocessor, nor are they converted

– *A macro may evaluate its arguments more than once*

Unexpected behavior may occur if an argument has side effects:

```
n = MAX(i++, j);
```

The same line after preprocessing:

```
n = ((i++) > (j)) ? (i++) : (j);
```

If *i* is larger than *j*, then *i* will be (*incorrectly*) *incremented twice* and *n* will be assigned an unexpected value

Parameterized Macros

- Errors caused by evaluating a macro argument more than once can be difficult to find, because a macro invocation looks the same as a function call
- To make matters worse, a macro may work properly most of the time, failing only for certain arguments that have side effects
- For self-protection, it is *a good idea to avoid side effects in arguments*

Parameterized Macros

- Parameterized macros can be used as patterns for segments of code that are often repeated

- A macro that makes it easier to display integers:

```
#define PRINT_INT(n) printf("%d\n", n)
```

- The preprocessor will turn the line

```
PRINT_INT(i/j);
```

into

```
printf("%d\n", i/j);
```