

Unix Filenames

- A file name consists of characters
- Almost *any character* is a valid file name character, including
 - all the punctuation and digits
 - the one exception is the / (slash) character (*the only invalid character in a Unix file name*)
 - the following characters are *not recommended* to be used in a Unix file name
 - ? * [] " ' () & ; ; !
 - the following are *not recommended* to be *used as the first character* in a Unix file name
 - - ~
 - control characters are also *allowed*, but are *not recommended*
- UPPER and lower case letters are different
 - A.txt and a.txt are different files

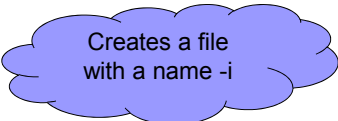
Unix Filenames

- Unix does not enforce file name extensions
 - The following are all legal Unix file names
 - a
 - a.
 - .a
 - ...
 - a.b.c
- Remember any name beginnings with a dot is hidden
 - ls cannot see them, but ls -a can
- . and .. are reserved for current and parent directories

Unix Filenames

- Executable files usually have no extensions
- Distinguishing executable files from data files might be difficult
- “file” command can help in this regard
 - Use: `file filename`
 - Result: print the type of the file
 - Example: `obelix[1] > file ~/.cshrc`
 `.cshrc: executable c-shell script`
- Filenames and pathnames have limits on lengths
 - Typically, can consist of 255 characters
 - these are pretty long (much better than the late MS-DOS days and the 8.3 filenames)

Fixing Filename Mistakes

- It is very easy to get the wrong stuff into file names
 - Say you accidentally typed
 `obelix[3] > cp myfile -i` • • • 
 - If you attempt to delete it as follow:
 `obelix[4] > rm -i`
 - The shell thinks `-i` is an option, not a file
 - Getting rid of these files can be painful
- There is an easy way to fix this...
 - You simply type
 `obelix[5] > rm -- -i`
 - Many commands use “--” to say there are no more options

Filename Wildcarding

- Wildcarding is the use of “*special*” characters to represent or match a sequence of other characters
 - A short sequence of characters can match
 - a long file name
 - many file names
- Wildcard characters include:
 - * matches a sequence of *zero* or more characters
 - Example: `a*.c*` matches for example `abc.c`, `apple.cpp`
 - ? matches any *single* character
 - Example: `a?.c` matches for example `ab.c`, `ax.c`, but *not* `abc.c` or `a.c`
 - [...] matches any *single* character between the braces
 - Example: `b[ae]t` matches exactly `bat`, `bet`, or `bit`, *not* `baet` or `bt`

Within [...], a pair of characters separated by “-” matches any character lexically between the two

 - Example: `[a-z]*` matches any file starts by a lower-case letter

Filename Wildcarding

- Wildcard sequences can be combined
 - obelix[6] > `mv a*.ch cfiles/`
 - `mv` all files beginning with `a` and ending with `.c` or `.h` into the directory `cfiles`
 - obelix[7] > `ls [abc]*.?`
 - list files whose name begins with `a`, `b`, or `c` and ends with `.` (dot) followed by a single character
- Wildcards do *not* cross “/” boundaries
 - Example: `cs2211*p` *does not* match `cs2211/tmp`
 - Example: `cs2211/*p` *does* match `cs2211/tmp`
- *Wildcards are expanded by the shell*
 - *The program just sees the list of files matched*

Filename Wildcarding

■ Matching the dot hidden files

- A dot (.) at the beginning of a filename must be matched explicitly
- Example: `ls .c*`



What does it mean?

- Note that, `ls ?c*` will **not** match `.c`

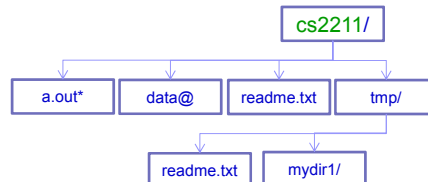
Filename Wildcarding

■ More advanced examples:

- What does the following command do?
`obelix[8] > ls -l /bin/*[-+]*`
- What the differences between
`obelix[9] > ls`
`obelix[10] > ls *`
`obelix[11] > ls */*`

List all non-hidden files or directories in the /bin directory that include at least a "-" or "+" in their name

Remember: Wildcards are expanded by the shell and the program just sees the list of files matched



Filename Wildcarding

- What about this?

```
obelix[12] > mv *.bat *.bit
```

If “*.bat” and “*.bit” **match exactly two file names**, then the “mv” command will rename the first file to the second one (the second file will be overridden)

If both “*.bat” and “*.bit” **match less than two file names**, you will get an error message saying “mv: Insufficient arguments”

If “*.bat” and “*.bit” **match more than two file names and the last one is not a directory**, you will get an error message saying that “mv: Target zzz.bit must be a directory”, where zzz.bit is the last file name in the list.

If “*.bat” and “*.bit” **match more than two file names and the last one is a directory**, all files will be moved to that directory

Unix Quoting

■ Double Quotes: “...”

- Putting text between double quotes

- **stops** the interpretation of *some* shell special characters (*whitespaces* and *wildcards*)
- **forces** the shell to deal with whatever between the two quotes as one token

- Examples:

```
obelix[13] > echo Here are some words
```

```
Here are some words
```

```
obelix[14] > echo "Here are some words"
```

```
Here are some words
```

```
obelix[15] > mkdir "A directory name with spaces?"
```

```
obelix[16] > ls A*
```

```
A directory name with spaces?/
```

- **Question:** what will happen if you execute:
mkdir A directory name with spaces?

How many arguments do we have in each of the above commands?

Unix Quoting

- Single Quotes ‘...’
 - Putting text between single quotes **stops** the interpretation of *even more* special characters
 - Stop variable expansion (\$HOME, etc.)
 - Examples:


```
obelix[17] > echo "Welcome $HOME"
Welcome /faculty/elsakka
obelix[18] > echo 'Welcome $HOME'
Welcome $HOME
```

Unix Quoting

- Back Quotes `...`
 - Putting text between back quotes *execute the text as a command and return the results*
 - Note the difference:
 - single quote (')
 - back quote (`)
 - Examples:


```
obelix[19] > cd
obelix[20] > echo "pwd"
pwd
obelix[21] > echo 'pwd'
pwd
obelix[22] > echo `pwd`
/faculty/elsakka
```

Unix Quoting

■ Backslash \

- ☐ *quotes* the next character
- ☐ Escapes *all* of the shell special characters


```
obelix[23] > mkdir Dir\ name\ with\ spaces\$*
```

```
obelix[24] > ls Dir\ *
```

```
Dir name with spaces$*/
```

- ☐ Use backslash to escape a newline character

```
obelix[25]> echo "This is a long line and\
we want to continue on the next one"
This is a long line and
we want to continue on the next one
```

Must be followed
by Enter

- ☐ Use backslash to escape other shell special characters

```
obelix[26] > echo \"This is Mahmoud's book\"
\"This is Mahmoud's book\"
obelix[27] > echo \"This is Mahmoud's book\"
\"This is Mahmoud's book\"
```

How many
arguments do we
have in each of
the above
commands?

Unix Quoting

■ Exercise (*enjoy it at home ☺*)

- What is the difference between the following commands:
(*Try to expect the result before hitting enter*)

- | | |
|---|---|
| <input type="checkbox"/> echo ls \$HOME | <input type="checkbox"/> echo \"ls \$HOME\" |
| <input type="checkbox"/> echo "ls \$HOME" | <input type="checkbox"/> echo \'ls \$HOME\' |
| <input type="checkbox"/> echo 'ls \$HOME' | <input type="checkbox"/> echo \"ls \$HOME\" |
| <input type="checkbox"/> echo `ls \$HOME` | <input type="checkbox"/> echo \"ls \$HOME\" |
| <input type="checkbox"/> echo ls \ \$HOME | <input type="checkbox"/> echo \'ls \$HOME\' |
| <input type="checkbox"/> echo "ls \ \$HOME" | <input type="checkbox"/> echo \"ls \$HOME\" |
| <input type="checkbox"/> echo 'ls \ \$HOME' | <input type="checkbox"/> echo \'ls \$HOME\' |
| <input type="checkbox"/> echo `ls \ \$HOME` | <input type="checkbox"/> echo \"ls \$HOME\" |
| <input type="checkbox"/> echo \"ls \$HOME\" | <input type="checkbox"/> echo \'ls \$HOME\' |
| <input type="checkbox"/> echo \'ls \$HOME\' | <input type="checkbox"/> echo \"ls \$HOME\" |
| <input type="checkbox"/> echo \"ls \$HOME\" | <input type="checkbox"/> echo \'ls \$HOME\' |

Hard and Symbolic Links

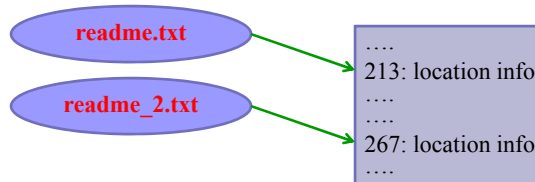
- Unix creates a single *index node* (*inode*) control structure for each file
- When a file is created, there is one link to the file's *inode*
- `ls -li` display the inode number for each file
- `ls -l` display the number of existing links per file
- Additional links can be added to a file using the command `ln` (**hard links**)
- Each hard link acts like a pointer to the actual file

Hard and Symbolic Links

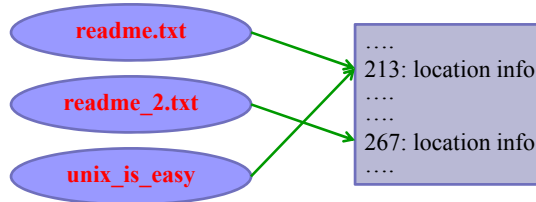
```

obelix[1] > ls -li
213 readme.txt          267 readme_2.txt
obelix[2] > ls -li
-rw-r--r-- 1 elsakka faculty 3456 Sept 10 23:34 readme.txt
-rw-r--r-- 1 elsakka faculty 876 Sept 8 21:02 readme_2.txt

```



Hard and Symbolic Links



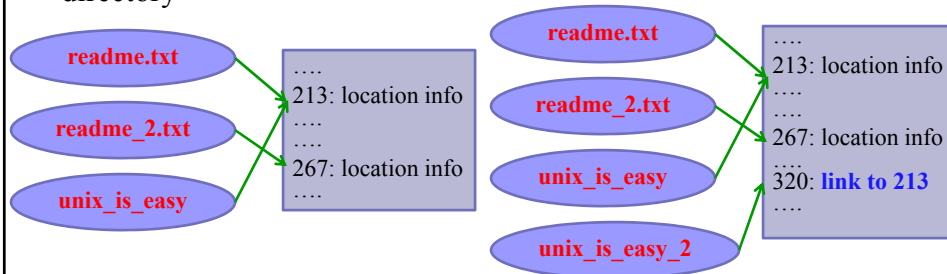
```

obelix[3] > ln readme.txt unix_is_easy
obelix[4] > ls -i
213 readme.txt          267 readme_2.txt 213 unix_is_easy
obelix[5] > ls -l
-rw-r--r--  2 elsakka faculty 3456 Sept 10 23:34 readme.txt
-rw-r--r--  1 elsakka faculty  876 Sept  8 21:02 readme_2.txt
-rw-r--r--  2 elsakka faculty 3456 Sept 10 23:34 unix_is_easy
  
```

- There is only one copy of the file contents on the hard disk, but has two distinct names!

Hard and Symbolic Links

- A symbolic link “ln -s” is an *indirect* pointer to another file or directory



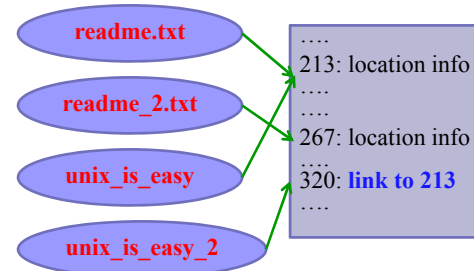
```

obelix[6] > ln -s readme.txt unix_is_easy_2
obelix[4] > ls -i
213 readme.txt          267 readme_2.txt 213 unix_is_easy 320 unix_is_easy_2
obelix[5] > ls -l
-rw-r--r--  2 elsakka faculty 3456 Sept 10 23:34 readme.txt
-rw-r--r--  1 elsakka faculty  876 Sept  8 21:02 readme_2.txt
-rw-r--r--  2 elsakka faculty 3456 Sept 10 23:34 unix_is_easy
lrwxrwxrwx  1 elsakka faculty  10 Sept 10 23:36 unix_is_easy_2 -> reamme.txt
  
```

Hard and Symbolic Links

■ What will happen if I delete:

- ☐ reame.txt
- ☐ unix_is_easy_2
- ☐ reame.txt and unix_is_easy_2
- ☐ reame.txt and unix_is_easy



```

obelix[4] > ls -l
213 readme.txt          267 readme_2.txt  213 unix_is_easy  320 unix_is_easy_2
obelix[5] > ls -l
-rw-r--r--    2  elsakka  faculty 3456  Sept 10 23:34  readme.txt
-rw-r--r--    1  elsakka  faculty  876  Sept  8 21:02  readme_2.txt
-rw-r--r--    2  elsakka  faculty 3456  Sept 10 23:34  unix_is_easy
lrwxrwxrwx    1  elsakka  faculty  10   Sept 10 23:36  unix_is_easy_2 -> reamme.txt

```

Hard and Symbolic Links

```

obelix[1] > cd
obelix[2] > ln -s /usr/local/bin bin.bin
obelix[3] > ls -l
lrwxrwxrwx    1  elsakka  faculty 14   Sept 11 01:34  bin.bin -> /usr/local/bin
.....
obelix[4] > cd bin.bin
obelix[5] > pwd
/usr/local/bin

```

Hard and Symbolic Links

- Hard links will have the same authority to a file
 - Removing any one of them, but one, *will not* remove the contents of the file
 - Removing *all* of the hard links *will* remove the contents of the file
- A symbolic link is just an entry pointing to the real name
 - Removing the symbolic link *does not* affect the file
 - Removing the original file will remove the contents of the file and leave the symbolic link pointing to unknown file
- Any user can create symbolic links to directories or files
- Any user can create hard links to files
- Only super users can create hard links to directories
- A hard link must point to a file in the same Unix filesystem
- Before Windows *vista*, there was no hard link under Windows (Only soft link, a.k.a. symbolic link)
- Even today, you can not create a hard link under windows using GUI (you have to use the *Windows* text command *mklink /H* to do so)

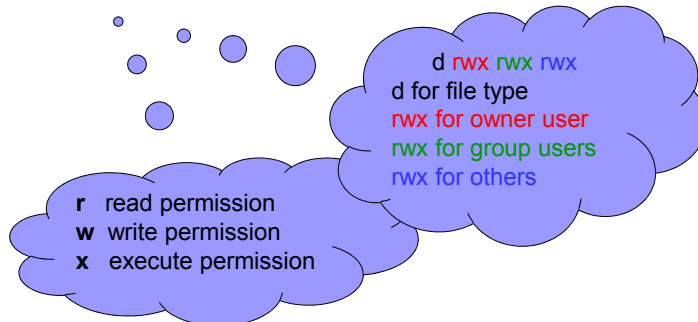
File Permissions

- Unix divides the set of all users in a system into three categories:
 - *owner user*
 - The owner of the file
 - *group users*
 - Users can be in more than one group
 - Each file is assigned to *one* and *only one* group
 - Most of you are in the group *2ndyr*
 - Used for easier access control (administration)
 - Normally only the *superuser* can set up groups
 - *other users*
 - Everyone else
- Unix allows you to give *distinct permission privilege* to each category
- The union of *owner user*, *group users*, and *other users* is a universal set
- The intersection between any two groups is an empty set

File Permissions

- Permissions can be viewed with the `ls -l` command

```
obelix[1] > ls -l
total 1247
-rw----- 1 elsakka faculty 1117 Sep 14 15:49 bad.cpp
drwx--x--x 2 elsakka faculty 2048 Oct 11 10:13 prog/
drwxr-xr-x 2 elsakka faculty 512 Sep 14 23:18 test2/
-rw-r-xr-- 1 elsakka faculty 2081 Jul 23 15:49 postal_code.txt
```



File Permissions

- Permissions are changed with the `chmod` command

- There are two syntaxes you can use:

`chmod DDD file [file ...]`

- `DDD` are 3 *octal* digits representing bits of protection
- `rw` `rw` `rw` can be thought of as `111 111 111` in binary
- This is the “numeric” method
- It can only perform “absolute” permission assignments, e.g.,

- `chmod 644 file`

```
rw-  r--  r--
110 100 100
 6   4   4
```

File Permissions

`chmod [ugoa][+|=][rwx] file [...]`

- This is the “symbolic” method
- It can perform “relative” permission assignments, e.g.,
 - `chmod u+rw file` gives the owner user Read, Write, and eXecute
 - `chmod g+rx file` gives the group users Read and eXecute
 - `chmod o-rwx file` removes from others Read, Write, and eXecute
 - `chmod a+x file` gives All eXecute permission
- It can perform “absolute” permission assignments, e.g.,
 - `chmod g=r file` gives group users Read permission and makes sure it has nothing else
- “relative” and “absolute” permission assignments can be appended with commas (*with no white space before or after the comma*)
 - `chmod u=rwx,g-w,o-rwx file`

The Umask Command

- `umask` sets the *default* permissions for any file you will create
- The meaning of `umask` (*un-mask*) value is opposite to that of the `chmod` command
 - tells which permissions will *not* be given (i.e., *un-masked*)
 - `umask 077` means do *not* let anyone but the owner user to do anything with my files by default
- Generally, set `umask` once in your `.cshrc` file and never set it again

Directory Permissions

- Directory permissions are *different* from the file permissions
 - Listing the contents of the directory requires the *read* permission to be set
 - Creating files in the directory requires the *write* permission to be set
 - *Accessing* files (*if* and *only if* you know its name) in the directory requires the *execute* permission to be set
- *drwx--x---* means it is a directory and it is
 - fully accessible to owner user
 - accessible by name (if you know the name) to group users
 - not accessible to others (including all subdirectories, if any)

Directory Permissions

- The *-R* option in *chmod* is useful when working with directories
 - It *recursively* changes the mode for each *chmod* operand that is a directory
 - All sub-directories and files receive those permissions
- *chmod -R a+r dir*
 - gives everyone read permission to
 - *dir* directory and
 - all sub-directories and files under *dir*
- *chmod -R a+x dir*
 - gives the executable access to allow everyone to
 - access *dir* directory and
 - makes all sub-directories and files under *dir* executable