

Unix Processes and Job Control

*Computer Science Department
CS2211a: Software Tools & Systems Programming
Fall 2013
Instructor: Mahmoud R. El-Sakka
Office: MC-419
Email: elsakka@csd.uwo.ca
Phone: 519-661-2111 x86996*

1


Topic 06: Unix Processes and Job Control

Introduction

- Within Unix, every **object** is represented by either:
 - **A file** (can be an input source or an output target)
 - **A process** (a program that is executing)
- A **process** is a **program** that is started by either
 - a **user**
 - the **system**
- **Processes** need to **share** the system's **resources**, including processors, memory, I/O devices, network connections,
- All **processes** are managed by the **kernel** through a **scheduling service** called **scheduler**

Introduction

- In a *single-processor* computer,
 - only one process at a time can be running
- Any **process** can have *one status* (*out of three*) at a time
 - **Running**
 - For a short slice of time only
 - **Ready to run**
 - A ready process means
 - it could use the CPU immediately (if it is allowed to do so)
 - The scheduler decides which of the ready processes to run next
 - **Blocked or sleeping**
 - Waiting for an event to occur, e.g., I/O or signal



© Mahmoud R. El-Sakka

3

CS 2211: Software Tools & Systems Programming

Introduction

- When a process is created, the kernel assigns it a unique identification number called **Process ID**, or simple **PID**
- To *keep track of all processes in the system*,
 - the kernel maintains a **process table**,
 - indexed by the **PID**,
 - containing one entry for each process
- The **scheduler**
 - **Maintains** a list of all the processes waiting to execute
 - **Chooses** one process at a time, and gives it a chance to run for a short time called **time slice** (*typically 10 milliseconds*)
 - **Saves data** related to the interrupted processes so that they can be resumed later
 - *Which data are we talking about?*

Similar to the idea of **inodes** in case of files, or **S.I.N.** in our daily life

© Mahmoud R. El-Sakka

4

CS 2211: Software Tools & Systems Programming

System Calls

- When a process needs the kernel to perform a service,
 - it sends a request (a **system call**) which will **create a new process**
- **System calls** provide an **essential interface** between
 - a process and
 - the operating system.

*The topic of System calls is important as it is essential in understanding many other subjects, including **operating systems** and **computer networks***

- Every created process (**child process**) is generated by another process (**parent process**)—*there is only one exception to this fact (we will talk about it later)*

System Calls

- Unix systems typically have between 200-300 **system calls**
- **System calls** may include
 - **File related calls** (e.g., opening, closing, reading from, or writing to a file)
 - **Information related calls** (e.g., get or set, current date/time)
 - **Communication related calls** (e.g., interprocess communication)
 - **Process related calls** (e.g., creating, executing, or stopping a process)
- Some of the **system calls** that **coordinate processes** include:
 - **fork**: **creates a copy (child)** of the current process (**parent**)
 - **exec**: **changes the program** that a process is running
 - **wait**: **forces a** process to **pause** until another process finishing its execution
 - **exit**: **terminates** a process

Internal vs External Commands

- There are two types of commands
 - *Internal commands*:
 - Interpreted directly by the shell (*no* need to create a new process)
 - *External commands*:
 - Required the shell to run a separate program (need to create a new process)

How to Run a Program or an External Command

- The shell will use *fork* system call to create a *duplicate of the currently running program*
 - The *duplicate* (child process) and the *original* (parent process)
 - proceed from the point of the fork with exactly the same data.
 - The only difference is the *return value from the fork call*.
fork returns
 - the child PID to the parent process and
 - 0 to the child process
 - By checking the return value of the *fork* call, each process can realize if it is a child or parent process

Process A

At time t

Process A1

Process A

At time $t + \Delta$

How to Run a Program or an External Command

- The child process
 - Identifies itself from the return value of the *fork* (i.e., the 0 value)
 - Uses an *exec* system call to change itself from a process running the shell into a process running the external command
- At the same time, the parent process
 - Identifies itself from the return value of the *fork* (i.e., the positive value)
 - Uses the *wait* system call to pause itself until the child is finishing the execution

How to Run a Program or an External Command

- When the child process finishes execution, it
 - Uses the *exit* system call to
 - Send a *signal* to the parent process indicating that it is done
 - Go to a *Zombie* state waiting for the terminating status to be accepted
 - Once a parent process receives a wake up *signal* from a child
 - It checks the outcome of the short life of its child
 - The kernel
 - de-allocates the child used resources
 - Removes the child PID from the process table (removing the last leftover from the that child)
 - Now the child is *declared dead/terminated* (or actually *killed*)

The very First Process

- Every Unix system has a process that is the *parent of all other processes* in the system
- In Unix, toward the end of the boot procedure, the kernel creates a special process *by hand without forking*
 - This process is given a PID of 0 (*process#0*)
 - *Process#0* initializes many data structures needed by the kernel
 - At the end, it forks *process#1* and then somehow disappeared (*terminated itself*)
 - *Process#1* (*a.k.a. the init process*) carries out the rest of the steps that are necessary to set up the kernel and finishes the boot process (*in doing so, many processes are created by init*)
 - Consequently, *init* becomes *the only living ancestor* of all other processes in the system

The very First Process

- *Process#1 (init)*
 - Is the first process in the process table
 - Stays there until the system is shut down
 - Is *the only living ancestor* of all other processes in the system
- If a parent process unexpectedly died,
 - The child will keep executing,
 - The child process is considered to be an *orphan*
 - *Process#1* automatically *adopts* all *orphan processes*, hence it
 - *receives the control* from the child process when finishing execution
 - *Informs* the kernel to *de-allocates the child used resources* and *kill it*

Running a Program in the Background

- In all examples seen so far, commands were run in the *foreground*
 - The shell waits for the command to finish before giving another prompt to allow you to continue
- When running a command in the *background*, you do not have to wait for the command to finish before starting another command
 - Useful when running a command that needs a long time
 - The window will be free so you can use it for other work

Running a Program in the Background

- One way to run a command in the background is to type ampersand, i.e., **&**, at the end of the command line
 - The shell will
 - display the job number (*job ID*) that identifies the command
 - display the Process ID (*PID*) number for each running command in the background
 - give you another prompt
- *A job* refers to *all the processes* that are necessary to execute an *entire command line*